

Assignment:-

Applying SGD on Boston House Prices

1. About the dataset Title: Boston House Prices dataset. Link: <http://archive.ics.uci.edu/ml/datasets/Housing> (<http://archive.ics.uci.edu/ml/datasets/Housing>).

Relevant Information: This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University. The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics..', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter. The Boston house-price data has been used in many machine learning papers that address regression problems.

Data includes: Number of Instances: 506 Number of Attributes: 13 numeric/categorical predictive *Median Value (attribute 14) is usually the target

Attribute Information: CRIM per capita crime rate by town ZN proportion of residential land zoned for lots over 25,000 sq.ft. INDUS proportion of non-retail business acres per town CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise) NOX nitric oxides concentration (parts per 10 million) RM average number of rooms per dwelling AGE proportion of owner-occupied units built prior to 1940 DIS weighted distances to five Boston employment centres RAD index of accessibility to radial highways TAX full-value property-tax rate per tenthousand dollar PTRATIO pupil-teacher ratio by town B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town LSTAT % lower status of the population MEDV Median value of owner-occupied homes in \$1000's

1. Objective:-

In Boston House Price dataset where we need to predict house price for a given set of attributes. So, this is a regression problem and we will use linear-regression to predict the house prices. Apart from that we will implement linear regression using both gradient descent optimizer and stochastic gradient descent(SGD) optimizer and will compare their performance.

In [1]:

```
# Loading required libraries
%matplotlib inline
import numpy as np
import pandas as pd
import scipy.stats as stats
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn
import string
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

2. Loading boston data set from sklearn.

In [2]:

```
# Loading boston dataset from sklearn
from sklearn.datasets import load_boston
boston=load_boston()

#Knowing datapoints and feature names
print('This dataset contains data about {} homes and each containing {} features about them')

#The boston variable itself is a dictionary, so we can check for its keys using the snippet
print(boston.keys())

#Knowing column names in our dataset
print(boston.feature_names)
```

This dataset contains data about 506 homes and each containing 13 features about them

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

CRIM: Per capita crime rate by town

ZN: Proportion of residential land zoned for lots over 25,000 sq. ft

INDUS: Proportion of non-retail business acres per town

CHAS: Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)

NOX: Nitric oxide concentration (parts per 10 million)

RM: Average number of rooms per dwelling

AGE: Proportion of owner-occupied units built prior to 1940

DIS: Weighted distances to five Boston employment centers

RAD: Index of accessibility to radial highways

TAX: Full-value property tax rate per \$10,000

PTRATIO: Pupil-teacher ratio by town

B: $1000(B_k - 0.63)^2$, where B_k is the proportion of [people of African American descent] by town

LSTAT: Percentage of lower status of the population

MEDV: Median value of owner-occupied homes in \$1000s

In [3]:

```
print(boston.target)
```

```
[24.  21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15.  18.9 21.7 20.4
 18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
 18.4 21.  12.7 14.5 13.2 13.1 13.5 18.9 20.  21.  24.7 30.8 34.9 26.6
 25.3 24.7 21.2 19.3 20.  16.6 14.4 19.4 19.7 20.5 25.  23.4 18.9 35.4
 24.7 31.6 23.3 19.6 18.7 16.  22.2 25.  33.  23.5 19.4 22.  17.4 20.9
 24.2 21.7 22.8 23.4 24.1 21.4 20.  20.8 21.2 20.3 28.  23.9 24.8 22.9
 23.9 26.6 22.5 22.2 23.6 28.7 22.6 22.  22.9 25.  20.6 28.4 21.4 38.7
 43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
 18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22.  20.3 20.5 17.3 18.8 21.4
 15.7 16.2 18.  14.3 19.2 19.6 23.  18.4 15.6 18.1 17.4 17.1 13.3 17.8
 14.  14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
 17.  15.6 13.1 41.3 24.3 23.3 27.  50.  50.  50.  22.7 25.  50.  23.8
 23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
 37.9 32.5 26.4 29.6 50.  32.  29.8 34.9 37.  30.5 36.4 31.1 29.1 50.
 33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50.  22.6 24.4 22.5 24.4 20.
 21.7 19.3 22.4 28.1 23.7 25.  23.3 28.7 21.5 23.  26.7 21.7 27.5 30.1
 44.8 50.  37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29.  24.  25.1 31.5
 23.7 23.3 22.  20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
 29.6 42.8 21.9 20.9 44.  50.  36.  30.1 33.8 43.1 48.8 31.  36.5 22.8
 30.7 50.  43.5 20.7 21.1 25.2 24.4 35.2 32.4 32.  33.2 33.1 29.1 35.1
 45.4 35.4 46.  50.  32.2 22.  20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
 21.7 28.6 27.1 20.3 22.5 29.  24.8 22.  26.4 33.1 36.1 28.4 33.4 28.2
 22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21.  23.8 23.1
 20.4 18.5 25.  24.6 23.  22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
 19.5 18.5 20.6 19.  18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
 22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25.  19.9 20.8 16.8
 21.9 27.5 21.9 23.1 50.  50.  50.  50.  50.  13.8 13.8 15.  13.9 13.3
 13.1 10.2 10.4 10.9 11.3 12.3  8.8  7.2 10.5  7.4 10.2 11.5 15.1 23.2
  9.7 13.8 12.7 13.1 12.5  8.5  5.  6.3  5.6  7.2 12.1  8.3  8.5  5.
 11.9 27.9 17.2 27.5 15.  17.2 17.9 16.3  7.  7.2  7.5 10.4  8.8  8.4
 16.7 14.2 20.8 13.4 11.7  8.3 10.2 10.9 11.  9.5 14.5 14.1 16.1 14.3
 11.7 13.4  9.6  8.7  8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
 14.1 13.  13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20.  16.4 17.7
 19.5 20.2 21.4 19.9 19.  19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
 16.7 12.  14.6 21.4 23.  23.7 25.  21.8 20.6 21.2 19.1 20.6 15.2  7.
  8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
 22.  11.9]
```

In [4]:

```
#Create a DataFrame bos containing all the data to use in predicting Boston Housing prices.
bos = pd.DataFrame(boston.data)
print(bos.head())
```

	0	1	2	3	4	5	6	7	8	9	10	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	

	11	12
0	396.90	4.98
1	396.90	9.14
2	392.83	4.03
3	394.63	2.94
4	396.90	5.33

In [5]:

```
bos.columns = boston.feature_names
print(bos.head())
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

In [6]:

```
print(boston.target.shape)
```

(506,)

In [7]:

```
boston.target[:5]
```

Out[7]:

```
array([24. , 21.6, 34.7, 33.4, 36.2])
```

In [8]:

```
bos['PRICE'] = boston.target
print(bos.head())
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

In [9]:

```
print(bos.describe())
```

	CRIM	ZN	INDUS	CHAS	NOX	R
M \ count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
0						
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.28463
4						
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.70261
7						
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.56100
0						
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.88550
0						
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.20850
0						
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.62350
0						
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.78000
0						

	AGE	DIS	RAD	TAX	PTRATIO	
B \ count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
0						
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.67403
2						
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.29486
4						
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.32000
0						
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.37750
0						
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.44000
0						
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.22500
0						
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.90000
0						

	LSTAT	PRICE
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

In [10]:

```
X = bos.drop('PRICE', axis = 1)
Y = bos.PRICE
```

2.1 Standardizing Data

In [11]:

```
from sklearn.preprocessing import StandardScaler
X_scaler = StandardScaler().fit(X)
standardized_X = X_scaler.transform(X)
print(standardized_X.shape)
```

(506, 13)

3.Applying Stochastic Gradient Descent From Scratch

3.1 Calculating cost function

In [12]:

```
# The below function will compute the cost for each point:
def cal_cost(theta,X,y,m):
    #m = len(y)

    predictions = X.dot(theta)

    cost = (1/2*506) * np.sum(np.square(predictions-y))

    return cost
```

3.2 Determining Optimal Weights and cost

In [13]:

```
#The below method will compute the optimal weights and cost :
def stochastic_gradient_descent(X,y,theta,learning_rate=0.2,iterations=10):

    m = len(y) #length of the data set
    cost_value = np.zeros(iterations)

    for it in range(iterations):
        cost = 0.0

        for i in range(m):
            rand_ind = np.random.randint(0,m)
            X_i = X[rand_ind,:].reshape(1,X.shape[1])
            y_i = y[rand_ind].reshape(1,1)
            prediction = np.dot(X_i,theta)

            theta = theta - (2/m)*learning_rate*( X_i.T.dot((prediction - y_i)))
            cost += cal_cost(theta,X_i,y_i,m)

        cost_value[it] = cost

    return theta, cost_value
```

In [14]:

```
#Learning_rate
lr =0.2

#no. of iterations
n_iter = 100

theta = np.random.randn(14,1)

#adding the bias weight's features
X_b = np.c_[np.ones((len(standardized_X),1)),standardized_X]

# calling the sgd_function
theta_updated,cost_history = stochastic_gradient_descent(X_b,Y,theta,lr,n_iter)

print(cost_history)
```

```
[42984674.45981438 19065733.90753102 10100948.15109663 8472760.97784814
 5167627.84629632 3430928.84321833 3013691.84990894 3540639.58368588
 3088216.4269686 2745874.80156676 2540935.32050461 2396364.74808609
 3323220.08855247 3337823.65750375 3151228.09614569 2591115.6388071
 3282577.94878804 2381458.35051078 2667218.4559602 2635620.91858626
 3041193.08491147 2979550.96528399 2596305.11386917 2710421.34551427
 2450213.98256632 1860192.32301035 2885561.58217175 2579718.76019554
 2567065.01151585 2929302.56167601 2923785.25396308 2826613.60790927
 2819304.8996654 2566193.08040736 2513590.64232507 2981537.02644306
 2792195.58638274 2392780.84937154 3103699.88455786 2359831.42289333
 2740740.48878731 2725267.73357724 2615822.77890918 2600783.65765379
 2474455.15942701 3022727.71142724 2743010.22879668 2963275.24282101
 2472223.94112213 3086494.0306036 2938482.52444306 2379091.79684328
 2890373.52524421 2640192.34712254 2603094.61621942 2868723.85986166
 2757439.79199756 3283084.56033509 2682190.63531021 2172807.26742458
 2600522.76977445 3644139.43679201 3216155.34728603 2877625.24240584
 2599895.56115197 2680314.4561199 2674166.01654458 2849259.42555033
 2463001.62624044 2532173.25271 2812172.8526802 2296220.65092211
 2605900.55416047 3011939.46403327 3792499.93705481 3040413.43473542
 2818975.75947017 3212445.09303528 3025690.13502377 2964236.91601047
 2192285.56040966 2376422.74378408 3366382.60543423 2840758.48235684
 2407846.87668732 2590718.62530185 3031033.20320776 2623860.03788896
 2881006.83812535 2437364.26079213 2715281.96937225 2252623.71277724
 2369198.31902153 2872979.3472489 2449271.41140047 3105800.59590797
 2437438.44395425 2433990.72280406 2617958.81933454 2151702.83684743]
```

3.3 Calculating weights

In [15]:

```

print('Intercept Term(bias term) : {:.3f}\n'.format(theta_updated[0][0]))

print('*'*100)
print('Predicted Weights(without bias term :)')
weight_vector = theta_updated[1:]
print(weight_vector)

```

Intercept Term(bias term) : 22.357

Predicted Weights(without bias term :)

```

[[-1.02675933]
 [ 1.08694025]
 [ 0.15820468]
 [ 0.63007197]
 [-1.92812299]
 [ 2.79943447]
 [-0.05714188]
 [-2.89220411]
 [ 2.26386162]
 [-1.95011786]
 [-2.02916807]
 [ 0.81521694]
 [-3.59298243]]

```

3.4 Calculating Different metrics

In [16]:

```

y_predicted = X_b.dot(theta_updated)
y_predicted = y_predicted.ravel()

from sklearn import metrics

print('MAE:', metrics.mean_absolute_error(Y, y_predicted))
print('MSE:', metrics.mean_squared_error(Y, y_predicted))
print('RMSE:', np.sqrt(metrics.mean_squared_error(Y, y_predicted)))

```

MAE: 3.2196879117976898

MSE: 22.07329507367132

RMSE: 4.698222544076783

In [17]:

```

from scipy import stats
x = Y
y = y_predicted
slope, intercept, r_value, p_value, std_err = stats.linregress(x,y)

idx = ['slope', 'intercept', 'r_value', 'p_value', 'std_err']
data = np.array([slope, intercept, r_value, p_value, std_err])

print(pd.DataFrame(data= data, index= idx , columns =['values']))

regline = lambda S: 0.7466*x +6.0860359
S=np.array([x.min(),x.max()])

```

	values
slope	7.505434e-01
intercept	5.445072e+00
r_value	8.596936e-01
p_value	3.789703e-149
std_err	1.986444e-02

3.5 Plotting graph between Actual Prices vs Predicted prices

In [18]:


```

plt.scatter(Y, y_predicted,color=['red','blue'])
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Actual Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")

```

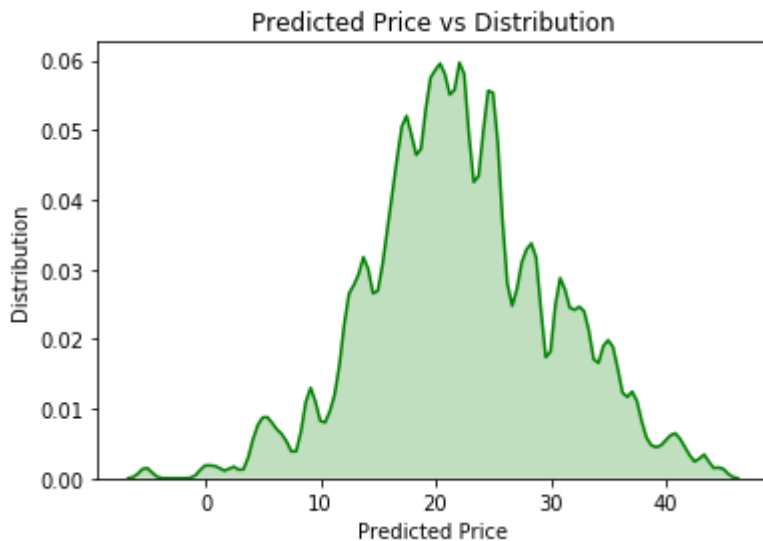
Out[18]:

Text(0.5,1,'Actual Prices vs Predicted prices: \$Y_i\$ vs \$\hat{Y}_i\$')



In [19]:

```
sns.kdeplot(y_predicted, bw = 0.5, color = "g", shade = True)
plt.xlabel("Predicted Price")
plt.ylabel("Distribution")
plt.title("Predicted Price vs Distribution")
plt.show()
```



4. Applying SGD Regressor from sklearn

In [20]:

```
#Computing the intercept and weight coefficients using the sklearn Library:
from sklearn import linear_model

clf =linear_model.SGDRegressor(n_iter = 100 , penalty=None ,eta0=0.2, loss='squared_loss' )
clf.fit(standardized_X , Y)
print("\n Accuracy of model using L1 Regularization",clf.score(standardized_X, Y))
#Predicting the target values for standardised data :
y_pred = clf.predict(standardized_X)
print('\n Number of coefficients',len(clf.coef_))
print("\n Intercept of the Best fit line : {}".format(clf.intercept_))
```

Accuracy of model using L1 Regularization 0.7271994230572881

Number of coefficients 13

Intercept of the Best fit line : [22.13122229]

4.1 Calculating intercepts and weights

In [21]:

```
# Intercept and weight vector Coffecient Calculation :
print('\n Intercept term :' , clf.intercept_ )
print('\n Weight vector :' , clf.coef_.reshape(-1 , 1))
```

Intercept term : [22.13122229]

Weight vector : [[-0.8294618]

```
[ 0.96874363]
[ 0.53636924]
[ 1.18495862]
[-2.28718941]
[ 2.13450772]
[ 0.09676401]
[-2.91377429]
[ 2.74079082]
[-1.82246344]
[-1.98559165]
[ 0.70719972]
[-3.91391879]]
```


4.2 Plotting a graph between Actual Prices vs Predicted prices

In [22]:

```
plt.scatter(Y, y_pred,color= ['red','blue'])
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Actual Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")
```


Out[22]:

Text(0.5,1,'Actual Prices vs Predicted prices: \$Y_i\$ vs \$\hat{Y}_i\$')



In [23]:

```
sns.kdeplot(y_pred, bw = 0.5, color = "g", shade = True)
plt.xlabel("Predicted Price")
plt.ylabel("Distribution")
plt.title("Predicted price vs Distribution")
plt.show()
```



4.3 Calculating different metrics

In [24]:

```
from sklearn import metrics

print('MAE:', metrics.mean_absolute_error(Y, y_pred))
print('MSE:', metrics.mean_squared_error(Y, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(Y, y_pred)))
```

MAE: 3.348244499466222
MSE: 23.029703624649628
RMSE: 4.798927341047125

In [25]:

```

from scipy import stats
x = Y
y = y_pred
slope, intercept, r_value, p_value, std_err = stats.linregress(x,y)

idx = ['slope', 'intercept', 'r_value', 'p_value', 'std_err']
data = np.array([slope, intercept, r_value, p_value, std_err])

print(pd.DataFrame(data= data, index= idx , columns =['values']))

regline = lambda S: 0.7466*x +6.0860359
S=np.array([x.min(),x.max()])

```

```

          values
slope      6.808384e-01
intercept  6.790022e+00
r_value    8.560335e-01
p_value    1.524802e-146
std_err    1.831309e-02

```

Comparing results of SGD from scratch and SGD Regressor

#SGD From Scratch

- 1) The Mean absolute error of SGD from scratch is 3.330089
- 2) The Root Mean Squared error of SGD from scratch is 4.713072

SGD Regressor from sklearn

- 1) The Mean absolute error of SGD Regressor from sklearn is 3.357264
- 2) The Root Mean Squared error of SGD Regressor from sklearn is 4.781386

Type of Algorithm	MSE	MAE	RMSE
SGD from scratch	22.073295	3.219687	4.698222

Type of Algorithm	MSE	MAE	RMSE
SGD Regressor from SKlearn	23.02970	3.348244	4.798927

Conclusion

Stochastic gradient descent also known as incremental gradient descent, is an iterative method for optimizing a differentiable objective function, a stochastic approximation of gradient descent optimization.

The gradient of the loss is estimated each sample at a time and the model is updated along the way with a decreasing strength schedule (aka learning rate).

The regularizer is a penalty added to the loss function that shrinks model parameters towards the zero vector using either the squared euclidean norm L2 or the absolute norm L1 or a combination of both (Elastic Net).

Learning rate. Learning rate is a decreasing function of time. Two forms that are commonly used are a linear function of time and a function that is inversely proportional to the time t

Steps Involved:-

- 1) Loading Boston dataset from sklearn
- 2) Standardizing Data
- 3) Applying Stochastic Gradient Descent From Scratch
- 4) calculated weights and intercepts values
- 5) Calculating Different metrics like mean squared error, mean absolute error, Root mean square error
- 6) Plotting graph between Actual Prices vs Predicted prices
- 7) Applying SGD Regressor from sklearn
- 8) Calculating intercepts and weights
- 9) Plotting a graph between Actual Prices vs Predicted prices
- 10) Calculating different metrics
- 11) Comparing results of SGD from scratch and SGD Regressor from sklearn

In []: